

AWS Web Application Infrastructure Architecture Summary

Project: Scalable Web Application Infrastructure

Technology: Terraform Infrastructure as Code

Cloud Provider: Amazon Web Services (AWS)

Document Version: 1.0

Date: November 2025

Repository: <https://github.com/sthakur1985/aws-web-app-deploy>

Executive Summary

This document provides a comprehensive overview of a production-ready, scalable web application infrastructure deployed on AWS using Terraform. The architecture implements a 3-tier design with security best practices, high availability, and multi-environment support.

Key Benefits

- **Scalability:** Auto Scaling Groups with configurable capacity
 - **Security:** Multi-layer security with private subnets and centralized IAM
 - **High Availability:** Multi-AZ deployment across 2 availability zones
 - **Cost Optimization:** Environment-specific resource sizing
 - **Compliance:** Security frameworks and audit trail capabilities
-

Page 1: Architecture Overview & Network Design

□ 3-Tier Architecture Design

The infrastructure follows a modern 3-tier architecture pattern:

Presentation Tier (Public Layer)

- **Application Load Balancer (ALB):** Traffic distribution and SSL termination
- **Public Subnets:** 2 subnets across 2 Availability Zones
- **Security:** HTTPS enforcement, security headers, CIDR-based access control
- **Features:** HTTP/2 support, connection draining, health checks

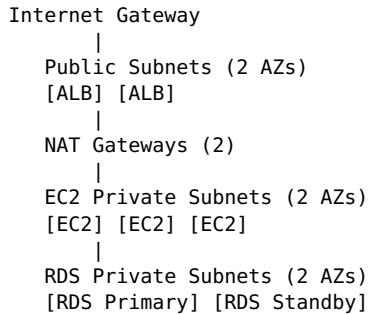
Application Tier (Private Layer)

- **Auto Scaling Group:** 1-10 EC2 instances based on environment
- **EC2 Private Subnets:** 2 subnets across 2 AZs with NAT Gateway access
- **Security:** No public IPs, security groups, centralized IAM roles
- **Features:** Automatic scaling, health monitoring, user data automation

Data Tier (Isolated Layer)

- **RDS MySQL/PostgreSQL:** Managed database service
- **RDS Private Subnets:** 2 subnets across 2 AZs, no internet access
- **Security:** Encryption at rest/transit, AWS Secrets Manager, network isolation
- **Features:** Multi-AZ deployment, automated backups, performance monitoring

Network Architecture



Subnet Design Strategy

- **CIDR Allocation:** /16 VPC with /24 subnets
- **Public Subnets:** 10.0.1.0/24, 10.0.2.0/24 (ALB placement)
- **EC2 Private:** 10.0.3.0/24, 10.0.4.0/24 (Application servers)
- **RDS Private:** 10.0.5.0/24, 10.0.6.0/24 (Database isolation)

Security Groups & Network ACLs

- **ALB Security Group:** HTTP/HTTPS from internet (0.0.0.0/0)
- **EC2 Security Group:** HTTP from ALB security group only
- **RDS Security Group:** MySQL/PostgreSQL port from EC2 security group only
- **Network Isolation:** RDS subnets have no internet routes

Security Architecture

Defense in Depth Strategy

1. **Network Layer:** VPC, subnets, security groups, NACLs
2. **Application Layer:** ALB security headers, HTTPS enforcement
3. **Data Layer:** Encryption, Secrets Manager, network isolation
4. **Identity Layer:** Centralized IAM, least privilege access
5. **Monitoring Layer:** VPC Flow Logs, CloudWatch alarms

Compliance & Governance

- **Encryption:** AES-256 for S3, RDS encryption at rest
- **Access Control:** Role-based access with IAM policies
- **Audit Trail:** CloudTrail logging, VPC Flow Logs
- **Secret Management:** AWS Secrets Manager for database credentials
- **Tagging Strategy:** Comprehensive resource tagging for cost allocation

Page 2: Service Components & Integration

▲ Content Delivery & DNS

CloudFront CDN

- **Global Distribution:** 400+ edge locations worldwide
- **Performance:** Intelligent caching, compression, HTTP/2
- **Security:** Origin Access Control (OAC), HTTPS enforcement
- **Cost Optimization:** Configurable price classes per environment

Route53 DNS Management

- **Hosted Zones:** Optional creation or existing zone integration
- **Alias Records:** ALB and CloudFront integration
- **Health Checks:** Automatic failover capabilities
- **Subdomain Strategy:** Configurable organization (app.domain.com, cdn.domain.com)

S3 Static Content

- **Security:** Public access blocked, OAC integration
- **Features:** Versioning enabled, server-side encryption
- **Integration:** CloudFront origin with bucket policies
- **Cost:** Lifecycle policies for storage optimization

Identity & Access Management

Centralized IAM Strategy

- **EC2 Role:** Systems Manager access, CloudWatch agent permissions
- **RDS Monitoring Role:** Enhanced monitoring capabilities (optional)
- **Instance Profiles:** Secure EC2 service access
- **Principle of Least Privilege:** Minimal required permissions

Security Best Practices

- **No Hardcoded Credentials:** AWS Secrets Manager integration
- **Service Roles:** Dedicated roles per AWS service
- **Managed Policies:** AWS-maintained security updates
- **Resource Tagging:** Complete metadata for governance

Monitoring & Observability

CloudWatch Integration

- **EC2 Monitoring:** CPU utilization, instance health
- **RDS Monitoring:** Database performance, storage capacity
- **ALB Monitoring:** Response times, error rates, unhealthy hosts
- **Custom Thresholds:** Environment-specific alert levels

Alerting Strategy

- **SNS Integration:** Email, SMS, webhook notifications
- **Escalation Policies:** Different thresholds per environment
- **Cost Control:** Optional monitoring features
- **Dashboard:** Centralized operational visibility

Deployment & Automation

Infrastructure as Code

- **Terraform Modules:** Reusable, parameterized components
- **Environment Separation:** Dev, staging, production configurations
- **State Management:** S3 backend with DynamoDB locking
- **Version Control:** Git-based infrastructure versioning

CI/CD Integration

- **GitHub Actions:** Automated deployment workflows
 - **Environment Promotion:** Consistent deployment across environments
 - **Secret Management:** Pipeline-based credential handling
 - **Testing:** Validation and dry-run capabilities
-

Page 3: Environment Strategy & Operational Considerations

Multi-Environment Architecture

Environment Differentiation

Component	Development	Staging	Production
EC2 Instances	t3.micro (1-2)	t3.small (1-3)	t3.medium (2-10)
RDS Configuration	Single-AZ, t3.micro	Multi-AZ, t3.small	Multi-AZ, t3.medium
Monitoring	Basic CloudWatch	Enhanced monitoring	Full monitoring + PI
Backup Retention	7 days	14 days	30 days
SSL Certificates	Optional	Required	Required
Flow Logs	Disabled	Enabled	Enabled
Cost Optimization	Aggressive	Balanced	Performance-focused

Environment Isolation Strategy

- **Separate AWS Accounts:** Complete isolation (recommended)
- **Separate VPCs:** Network-level isolation within account
- **Separate State Files:** Independent Terraform state management
- **Environment-Specific Variables:** Tailored configurations

Cost Optimization Strategy

Resource Sizing

- **Development:** Minimal resources, single AZ deployment
- **Staging:** Production-like but smaller scale
- **Production:** Performance and availability optimized

Cost Components (Monthly Estimates)

- **EC2 Instances:** \$8-50 per instance (t3.micro to t3.medium)
- **RDS:** \$15-100 (db.t3.micro to db.t3.medium, Multi-AZ)
- **ALB:** \$16 base + \$0.008 per LCU
- **NAT Gateways:** \$45 each (2 for HA)
- **CloudFront:** \$0.085 per GB + requests
- **S3:** \$0.023 per GB + requests

Cost Control Measures

- **Auto Scaling:** Automatic capacity adjustment
- **Reserved Instances:** Long-term cost reduction

- **Spot Instances:** Development environment optimization
- **Lifecycle Policies:** S3 storage class transitions

Operational Excellence

Deployment Process

1. **Code Review:** Terraform configuration validation
2. **Planning:** Dry-run execution with terraform plan
3. **Approval:** Manual approval for production changes
4. **Deployment:** Automated terraform apply
5. **Validation:** Post-deployment health checks

Disaster Recovery

- **Multi-AZ:** Automatic failover capabilities
- **Backup Strategy:** Automated RDS backups, S3 versioning
- **Infrastructure Recovery:** Terraform-based infrastructure recreation
- **Data Recovery:** Point-in-time recovery for RDS

Security Operations

- **Patch Management:** Systems Manager patch automation
- **Vulnerability Scanning:** Regular security assessments
- **Access Reviews:** Periodic IAM permission audits
- **Incident Response:** CloudWatch alarms and notification workflows

Performance Optimization

- **Auto Scaling Policies:** CPU and custom metric-based scaling
- **Database Optimization:** Performance Insights, parameter tuning
- **CDN Optimization:** CloudFront caching strategies
- **Application Monitoring:** Custom application metrics

Scalability & Future Enhancements

Horizontal Scaling

- **Auto Scaling Groups:** Automatic instance management
- **Load Balancer:** Traffic distribution across instances
- **Database:** Read replicas for read-heavy workloads
- **CDN:** Global edge location distribution

Future Enhancements

- **Container Integration:** ECS/EKS for containerized workloads
- **Serverless Components:** Lambda functions for specific tasks
- **Advanced Monitoring:** X-Ray tracing, custom dashboards
- **Security Enhancements:** WAF integration, GuardDuty monitoring

Conclusion

This AWS web application infrastructure provides a robust, scalable, and secure foundation for modern web applications. The architecture implements industry best practices for security, availability, and operational excellence while maintaining cost efficiency.

efficiency through environment-specific optimizations.

Key Success Factors: - Modular Terraform design for maintainability - Multi-environment strategy for safe deployments - Comprehensive security implementation - Automated monitoring and alerting - Cost-optimized resource allocation

Contact Information:

Technical Lead: Soumya Thakur (soumyathakur85@gmail.com)

Project Repository: AWS Web Application Infrastructure

Documentation: Complete module-level README files available